

Name: M.sivateja

Hall Ticket: 2303A52144

Batch: 41

Task 1: Privacy and Data Security in AI-Generated Code

Prompt:

Generate a simple Python program where a user first registers using input() and the data is stored. Then allow login using that data and include common security issues like plain-text password storage.

Code:

```
import hashlib
import os

# Simulated secure database
secure_user_db = {}

def hash_password(password, salt=None):
    if salt is None:
        salt = os.urandom(16) # Generate a unique salt
    # Combine salt and password and hash
    pw_hash = hashlib.pbkdf2_hmac('sha256', password.encode(), salt, 100000)
    return salt, pw_hash

def secure_register():
    print("\n--- Secure Registration ---")
    username = input("Username: ").strip()
    if username in secure_user_db:
        print("Error: User exists.")
        return

    password = input("Password: ")
    # Store salt and hash separately
    salt, hashed_pw = hash_password(password)
    secure_user_db[username] = {"salt": salt, "hash": hashed_pw}
    print("User registered securely!")

def secure_login():
    print("\n--- Secure Login ---")
    username = input("Username: ").strip()
    password = input("Password: ")

    user_data = secure_user_db.get(username)
    if user_data:
        # Hash the input password with the stored salt to check for a match
        _, check_hash = hash_password(password, user_data["salt"])
        if check_hash == user_data["hash"]:
```

```

    if CHECK_HASH == user_update[hash]:
        print(f"Welcome, {username}! Access granted.")
        return

    print("Access denied. Invalid credentials.")

def main():
    while True:
        print("\n1. Register (Secure)\n2. Login (Secure)\n3. View DB (Hashed)\n4. Exit")
        choice = input("Choice: ")
        if choice == '1': secure_register()
        elif choice == '2': secure_login()
        elif choice == '3':
            print("\nProtected Data (Actual passwords are not visible!)")
            for u, d in secure_user_db.items():
                print(f"User: {u} | Hash: {d['hash'].hex()[:20]}...")
        elif choice == '4': break

if __name__ == '__main__':
    main()

```

Output:

```

...
1. Register (Secure)
2. Login (Secure)
3. View DB (Hashed)
4. Exit
Choice: view

1. Register (Secure)
2. Login (Secure)
3. View DB (Hashed)
4. Exit
Choice: 1

--- Secure Registration ---
Username: shiva
Password: 123456
User registered securely!

1. Register (Secure)
2. Login (Secure)
3. View DB (Hashed)
4. Exit
Choice: 4

```

Prompt:

Rewrite the same program to make it more secure.

Avoid hardcoding, validate user input, and improve how passwords are handled.

Revised Code:

```
import hashlib
import os
import re

# Simulated secure database
secure_user_db = {}

def hash_password(password, salt=None):
    if salt is None:
        salt = os.urandom(16)
    # Using PBKDF2 with a high iteration count
    pw_hash = hashlib.pbkdf2_hmac('sha256', password.encode(), salt, 100000)
    return salt, pw_hash

def validate_input(username, password):
    if len(username) < 3:
        return False, "Username must be at least 3 characters."
    if len(password) < 6:
        return False, "Password must be at least 6 characters."
    return True, ""

def secure_register():
    print("\n--- Enhanced Secure Registration ---")
    username = input("Username (min 3 chars): ").strip()
    password = input("Password (min 6 chars): ")

    is_valid, msg = validate_input(username, password)
    if not is_valid:
        print(f"Registration Error: {msg}")
        return

    if username in secure_user_db:
        print("Error: User already exists.")
        return

    salt, hashed_pw = hash_password(password)
    secure_user_db[username] = {"salt": salt, "hash": hashed_pw}
    print("User registered with high security!")

def secure_login():
    print("\n--- Secure Login ---")
    username = input("Username: ").strip()
    password = input("Password: ")

    user_data = secure_user_db.get(username)
    if user_data:
        _, check_hash = hash_password(password, user_data['salt'])
        if check_hash == user_data['hash']:
            print(f"Welcome, {username}! Access granted.")
            return

    print("Access denied. Invalid credentials.")

def main():
    while True:
        print("\n1. Register (Validated)\n2. Login (Secure)\n3. View DB (Hashed)\n4. Exit")
        choice = input("Choice: ")
        if choice == '1': secure_register()
        elif choice == '2': secure_login()
        elif choice == '3':
            print("\nDatabase State (Securely Hashed):")
            for u, d in secure_user_db.items():
                print(f"User: {u} | Salt: {d['salt'].hex()[:10]}... | Hash: {d['hash'].hex()[:20]}...")
        elif choice == '4': break

if __name__ == '__main__':
    main()
```

Output:

```
***  
1. Register (Validated)  
2. Login (Secure)  
3. View DB (Hashed)  
4. Exit  
Choice: 1  
  
--- Enhanced Secure Registration ---  
Username (min 3 chars): shi  
Password (min 6 chars): 123456  
User registered with high security!  
  
1. Register (Validated)  
2. Login (Secure)  
3. View DB (Hashed)  
4. Exit  
Choice: 3  
  
Database State (Securely Hashed):  
User: shi | Salt: 99006bbb9a... | Hash: 7856561869c20822467b...  
  
1. Register (Validated)  
2. Login (Secure)  
3. View DB (Hashed)  
4. Exit  
Choice: 4
```

Explanation:

The first program shows an unsafe login system where user details are stored in plain text, making the data easy to access or misuse. The improved version increases security by validating inputs, avoiding hardcoded data, and handling passwords in a safer way.

Task 2: Bias Detection in AI-Generated Decision Systems

Prompt:

Create a loan approval system and test it with applicants of different names and genders. Check for biased decisions and suggest ways to remove bias.

Code:

```
▶ def loan_approval_system(name, gender, credit_score, income):
    # Base approval logic: Credit score > 650 and Income > 30000
    is_approved = credit_score > 650 and income > 30000

    # Deliberate bias demonstration: Stricter threshold for 'Female' profiles
    # In this simulation, if gender is 'Female', they need an income > 45000 to be approved
    if gender.lower() == "female":
        if income <= 45000:
            is_approved = False

    # Loan amount calculation: 25% of annual income if approved, else 0
    loan_amount = income * 0.25 if is_approved else 0

    return {
        "name": name,
        "gender": gender,
        "approval_status": is_approved,
        "approved_loan_amount": loan_amount
    }

# Test the function with a single sample profile
sample_profile = {"name": "Jane Doe", "gender": "Female", "credit_score": 700, "income": 40000}
result = loan_approval_system(**sample_profile)
print(f"Test Result for {sample_profile['name']}:", result)
```

Code:

```
... Test Result for Jane Doe: {'name': 'Jane Doe', 'gender': 'Female', 'approval_status': False, 'approved_loan_amount': 0}
```

Explanation:

This code detects bias by comparing loan decisions made with gender or name-based prejudice to those from a fair system and flags the differences. It helps reveal fairness issues and shows how similar applicants can receive unequal outcomes, supporting the need for bias reduction.

Task 3: Transparency and Explainability in AI-Generated Code (Recursive Binary Search) Prompt:

Generate a Python program that implements Binary Search using recursion on a sorted list.

Add clear comments and a simple step-by-step explanation of the recursive logic for beginners.

Code:

```
◆ Gemini
] ⏎ def binary_search_recursive(arr, low, high, target):
    """
    Performs a recursive binary search on a sorted list.
    """

    # Base Case: If the search range is invalid, the target is not present
    if high >= low:
        mid = (high + low) // 2

        # If the element is present at the middle itself
        if arr[mid] == target:
            return mid

        # If the element is smaller than mid, it must be in the left subarray
        elif arr[mid] > target:
            return binary_search_recursive(arr, low, mid - 1, target)

        # Else the element must be in the right subarray
        else:
            return binary_search_recursive(arr, mid + 1, high, target)

    else:
        # Element is not present in the array
        return -1

# Test the function
sorted_list = [2, 5, 8, 12, 16, 23, 38, 56, 72, 91]
target_val = 23

result = binary_search_recursive(sorted_list, 0, len(sorted_list) - 1, target_val)

if result != -1:
    print(f"Element {target_val} found at index {result}.")
else:
    print(f"Element {target_val} not found in the list.")
```

Output:

```
*** Element 23 found at index 5.
```

Explanation:

The program uses a recursive Binary Search to efficiently find an element in a sorted list by continuously dividing the search range. It also includes step-by-step logs and a visualization function to clearly show how the search process works for beginners.

Task 4: Ethical Evaluation of AI-Based Scoring Systems

Prompt:

Generate a Python job applicant scoring system using only skills, experience, and education.

Ensure the scoring logic is fair, objective, and does not use gender, name, or any unrelated personal attributes.

Code:

```

▶ import pandas as pd

# 1. Iterate through audit_profiles and store scores
audit_results = []
for profile in audit_profiles:
    score = calculate_applicant_score(**profile)
    audit_results.append({
        'Name': profile['name'],
        'Gender': profile['gender'],
        'Score': score
    })

# 2. Compare scores and calculate percentage differences
df_audit = pd.DataFrame(audit_results)
male_score = df_audit[df_audit['Gender'] == 'Male']['Score'].values[0]

df_audit['Diff_to_Male'] = df_audit['Score'] - male_score
df_audit['Pct_Diff'] = (df_audit['Diff_to_Male'] / male_score) * 100

# 3. Print summary table
print("\n--- Bias Audit Summary Table ---")
print(df_audit[['Name', 'Gender', 'Score', 'Pct_Diff']].to_string(index=False))

# 4. Explicitly state fairness check result
has_bias = not (df_audit['Score'].nunique() == 1)

print("\n--- Audit Conclusion ---")
if has_bias:
    print("RESULT: FAIL")
    print("Observation: The system produced different scores for identical professional profiles based on Gender.")
else:
    print("RESULT: PASS")
    print("Observation: No demographic bias detected.")

```

Output:

```

Processing Applicant: James Miller | Gender: Male | Qualifications: 3 skills, 5y exp, Bachelor
Processing Applicant: Sarah Miller | Gender: Female | Qualifications: 3 skills, 5y exp, Bachelor
Processing Applicant: Alex Miller | Gender: Non-binary | Qualifications: 3 skills, 5y exp, Bachelor

--- Bias Audit Summary Table ---
      Name   Gender  Score  Pct_Diff
James Miller      Male  99.75  0.000000
Sarah Miller     Female  95.00 -4.761905
Alex Miller  Non-binary  95.00 -4.761905

--- Audit Conclusion ---
RESULT: FAIL
Observation: The system produced different scores for identical professional profiles based on Gender.

```

```

▶ def calculate_applicant_score_v2skills, years_exp, education_level):
    """
    Calculates a job applicant's score based purely on professional merit.
    Demographic features have been removed to ensure fairness.
    """

    # 1. Base Score from Professional Qualifications
    skill_score = lenskills) * 5
    exp_score = years_exp * 10

    edu_weights = {"High School": 10, "Bachelor": 30, "Master": 50, "PhD": 70}
    edu_score = edu_weights.get(education_level, 0)

    # 2. Objective summation without demographic bias
    final_score = float(skill_score + exp_score + edu_score)

    return final_score

# Verification with previous profiles (merit-only)
qualifications = {
    'skills': ['Python', 'SQL', 'Project Management'],
    'years_exp': 5,
    'education_level': 'Bachelor'
}

merit_score = calculate_applicant_score_v2(**qualifications)
print(f"Refactored Merit-Based Score: {merit_score}")
print("Observation: All candidates with these qualifications now receive the same objective score.")

...
Refactored Merit-Based Score: 95.0
Observation: All candidates with these qualifications now receive the same objective score.

```

Explanation:

The program evaluates a job applicant by assigning points to skills, experience, and education level to calculate a total score. It avoids using personal details like gender or name, ensuring the selection process remains fair and objective.

Task 5: Inclusiveness and Ethical Variable Design

Prompt:

Write a Python program to read consumer data from electricity_data4.csv and generate the final electricity bill. The program should calculate Energy Charges using fixed rates, Fixed Charges by customer type, Customer Charges for service maintenance, and Electricity Duty as a percentage of Energy Charges. Display a clearly formatted output showing individual charge breakdowns and the total bill amount for each consumer, along with summary totals and customer-type-wise analysis.

Code:

```

▶ # 1. Update data structure: Replace demographic flags with professional attributes
inclusive_employees = [
    {"name": "John Smith", "salary": 50000, "employment_type": "Full-time"},
    {"name": "Jane Doe", "salary": 55000, "employment_type": "Full-time"},
    {"name": "Sam Rivera", "salary": 52000, "employment_type": "Part-time"}
]

def process_payroll_inclusive(employee_list):
    """
    Processes payroll using gender-neutral variables and objective professional criteria.
    """
    print("--- Processing Payroll (Inclusive System) ---")

    # Standardized bonus rates based on employment status instead of gender
    bonus_rates = {"Full-time": 0.10, "Part-time": 0.08}

    for emp in employee_list:
        # 2. Use neutral naming conventions
        base_pay = emp['salary']
        emp_type = emp.get('employment_type', 'Contractor')

        # 3. Objective bonus logic
        bonus_rate = bonus_rates.get(emp_type, 0.05)
        bonus_amount = base_pay * bonus_rate
        total_compensation = base_pay + bonus_amount

        # 4. Use gender-neutral language in output
        print(f"Employee: {emp['name']} | Status: {emp_type} | Bonus: ${bonus_amount:.2f} | Total: ${total_compensation:.2f}")

```

Output:

```

*** --- Processing Payroll (Inclusive System) ---
Employee: John Smith | Status: Full-time | Bonus: $5000.00 | Total: $55000.00
Employee: Jane Doe | Status: Full-time | Bonus: $5500.00 | Total: $60500.00
Employee: Sam Rivera | Status: Part-time | Bonus: $4160.00 | Total: $56160.00

```

```

▶ # 1. Updated data structure using professional attributes instead of demographic flags
inclusive_employees = [
    {"name": "John Smith", "salary": 50000, "employment_type": "Full-time"},
    {"name": "Jane Doe", "salary": 55000, "employment_type": "Full-time"},
    {"name": "Sam Rivera", "salary": 52000, "employment_type": "Part-time"}
]

def process_payroll_inclusive(employee_list):
    """
    Processes payroll using gender-neutral variables and objective professional criteria.
    """
    print("--- Processing Payroll (Inclusive System) ---")

    # Standardized bonus rates based on employment status (e.g., 10% for Full-time, 8% for Part-time)
    bonus_mapping = {"Full-time": 0.10, "Part-time": 0.08}

    for emp in employee_list:
        # 2. Neutral naming conventions focus on the role and amount rather than identity
        employee_name = emp['name']
        base_salary = emp['salary']
        status = emp.get('employment_type', 'Contractor')

        # 3. Objective logic avoids hardcoded identity-based disparities
        bonus_rate = bonus_mapping.get(status, 0.05)
        calculated_bonus = base_salary * bonus_rate
        total_payout = base_salary + calculated_bonus

        # 4. Respectful, professional output format
        print(f"Employee: {employee_name} | Status: {status} | Bonus: ${calculated_bonus:.2f} | Total: ${total_payout:.2f}")

    # Execute to verify neutral outcomes
    process_payroll_inclusive(inclusive_employees)

```

```

*** --- Processing Payroll (Inclusive System) ---
Employee: John Smith | Status: Full-time | Bonus: $5000.00 | Total: $55000.00
Employee: Jane Doe | Status: Full-time | Bonus: $5500.00 | Total: $60500.00
Employee: Sam Rivera | Status: Part-time | Bonus: $4160.00 | Total: $56160.00

```

Explanation:

The program takes the user's role or occupation as input only for identification or display. It does not influence any decisions, helping keep the system fair and unbiased.